

Planning and Automated Reasoning - Automated Reasoning

I term, Academic Year 2025-26

Project: Implementation of a solver for the union of the theories of equality, lists, and arrays

Assigned November 13, 2025, to be submitted by 23:59 of January 31st, 2026

Assignment

The project consists of the implementation of a prototype solver that determines the satisfiability of a set (or conjunction) of literals in the union of the quantifier-free fragments of three theories: T_E (theory of equality with free (or uninterpreted) symbols), T_cons (theory of non-empty possibly cyclic lists), and T_A (theory of arrays without extensionality).

Since satisfiability in T_cons reduces to satisfiability in T_E, and satisfiability in T_A reduces to satisfiability in T_E, it is not necessary to implement the Nelson-Oppen scheme for theory combination. On the other hand, it is not enough to test whether the input problem S is in T_E, T_cons, or T_A, and then only invoke the T_E-procedure, the T_cons-procedure, or the T_A-procedure, respectively, because the three theories may be *mixed* in the input S. The input problem S may involve all three theories, or any two of them, or only one of them.

The solver should test first whether there is any occurrence of the **store** symbol of T_A, and decompose the problem into subproblems S_1, S_2, ... (two subproblems for each occurrence of **store**) according to the T_A-procedure seen in class and described in Sect. 9.5 of reference [1]. Once all occurrences of **store**, if any, have been processed away, for each subproblem S_k, the solver processes all occurrences of **select** of S_k according to the T_A-procedure. Once all occurrences of **select**, if any, have been processed away in each S_k, the solver tests each S_k in order to see whether S_k has any occurrence of T_cons-symbols. If yes, it applies to S_k the T_cons-procedure seen in class and described in Sect. 9.4 of reference [1], if not, it applies to S_k the T_E-procedure seen in class and described in Sect. 9.3 of reference [1]. The T_cons-procedure builds the T_cons-axioms into the congruence closure (CC) algorithm. The T_E-procedure consists of the CC algorithm.

The heart of the solver is the CC algorithm on DAG's. The CC algorithm as described in Sect. 9.3 of reference [1] should be improved with a non-arbitrary choice of the representative of the union class in the **UNION** function: pick the one with the largest **ccpar** set (see page 761 top in reference [8] and page 423 top in reference [7]). The other two variants discussed informally in class, namely the forbidden list or forbidden set (see page 388 bottom in reference [7]) and a non-recursive version of the **FIND** function (coupled with a version of the **UNION** function that updates the **find** field of all the terms in the class whose representative is not chosen as representative of the union class) are optional. Since these features are heuristic in nature, it is not guaranteed that they are helpful. They can be implemented as options, making it possible to evaluate their impact.

The programming language should be a general-purpose portable programming language that can be compiled on Linux, such as Java, C++, C, Rust, Standard ML, Ocaml. Programming languages that are oriented towards certain areas, such as Go (API, microservices, IOT), or R (data science, statistics), or Mathematica, Maple, MathScheme (symbolic computation) are not to be used. All the more Python is not to be used, because the point of the project is to develop a prototype from

scratch, not to assemble existing code. Similarly, translating existing code to another language is not a solution. The program will have an interface (stdin/stdout or a simple GUI) that allows the user to submit an input set (typically contained in a file) and get the answer.

In order to test the program, use literal sets from sources including the following:

1. Examples and exercises from books and papers, beginning with those in the book list of the course and in the references below.
2. Sets of literals obtained from more general formulas (also from books or papers) as follows, e.g.:
 - If the formula contains free predicate symbols, the solver handles them by the transformation explained in class;
 - If the formula is not a conjunction of literals, the solver transforms it into DNF, and then works on each disjunct;
 - If the formula contains defined (constant or function) symbols from theories other than lists and arrays (e.g., arithmetical symbols), the solver handles them by replacing them with free symbols;
 - If the formula contains quantifiers, the solver drops them and treats the variables as free variables.

Clearly, the last two transformations do not preserve equisatisfiability, and they are intended only to get more inputs for the solver.

3. *Optional:* Synthetic sets of literals obtained from a generator to be implemented in addition to the solver.
4. *Optional:* Benchmarks from the **SMT-LIB** library: the benchmarks in the **QF-UF** class should fall in the scope of the project. QF stands for quantifier free and UF stands for undefined function symbols. Since all logics/theories in the SMT-LIB repository feature equality, **QF-UF** stands for the quantifier-free fragment of the theory of equality. The benchmarks of SMT-LIB are currently hosted on Zenodo at <https://zenodo.org/records/16740866>. These benchmarks are written in the SMT-LIB language that is far more complex than the logic used in class, books, and papers. Thus, in order to give these inputs to the program, one needs a parser that handles at least the subset of SMT-LIB used in the **QF-UF** benchmarks.

The project also requires to write a report (max 6 pages 11pt) presenting

- The implementation, emphasizing major choices (e.g., data structures, heuristics), with comments on their impact (e.g., on ease of implementation, performance) or any other information deemed significant;
- A summary of the results of the experiments in the form of one or more tables or plots, reporting data such as answer (SAT/UNSAT), the run time, the source of the problems, or other data deemed relevant;
- Some analysis of the experiments, such as comments about performance, impact of features, or any other remarks deemed interesting.

The usage of generative AI tools to produce the report or parts thereof is **not** allowed, because the point of the report is to stimulate original thinking and writing.

Hand-in

1. A compressed archive in .tgz¹ or .zip format (.rar is forbidden) to be sent to the instructor by e-mail. The archive should contain:
 - (a) The source code,
 - (b) The input files used in the experiments, each including a comment stating the source of the problem,
 - (c) The output files corresponding to the input files,
 - (d) A **README** file with instructions on how to execute the program,
 - (e) Either a portable Linux executable, or a **README** file which contains also instructions on how to compile/make the program.

The archive should be named FirstNameLastNameStudentId (i.e., NomeCognomeNumeroMatricola).

2. A double-sided print-out of the report to be put in the instructor's mailbox.

It is not required to include the report in the archive and doing it does not fulfill the requirement of submitting separately the double-sided print-out of the report.

¹ E.g., **tar czvf - FolderName > FolderName.tgz** compresses and **tar xzpvf FolderName.tgz** decompresses.

References

1. Aaron R. Bradley, Zohar Manna. *The Calculus of Computation. Decision Procedures with Applications to Verification*. Springer, 2007, ISBN 978-3-642-09347-0.
2. Daniel Kroening, Ofer Strichman: *Decision Procedures. An Algorithmic Point of View*, Springer, 2008, ISBN: 978-3-540-74104-6.
3. Alessandro Armando, Silvio Ranise, and Michaël Rusinowitch. A rewriting approach to satisfiability procedures. *Information and Computation* 183(2):140–164, 2003.
4. Alessandro Armando, Maria Paola Bonacina, Silvio Ranise, and Stephan Schulz. New results on rewrite-based satisfiability procedures. *ACM Transactions on Computational Logic* 10(1):129–179, 2009.
5. Leo Bachmair, Ashish Tiwari and Laurent Vigneron. Abstract congruence closure. *Journal of Automated Reasoning* 31(2):129–168, 2003.
6. Aaron R. Bradley, Zohar Manna, and Henny B. Sipma. What's decidable about arrays? In E. Allen Emerson and Kedar S. Namjoshi (Eds.), *Proceedings of the Seventh International*

Conference on Verification, Model Checking and Abstract Interpretation (VMCAI), Lecture Notes in Computer Science 3055:427–442, Springer, 2006.

7. David L. Detlef, Greg Nelson and James B. Saxe. Simplify: a theorem prover for program checking. *Journal of the ACM* 52(3):365–473, 2005.
8. Peter J. Downey, Ravi Sethi, and Robert Endre Tarjan. Variations on the common subexpression problem. *Journal of the ACM* 27(4):758–771, 1980.
9. Dexter Kozen. Complexity of finitely presented algebras. Technical Report TR-76-294, Department of Computer Science, Cornell University, 1976.
10. Greg Nelson and Derek C. Oppen. Fast decision procedures based on congruence closure. *Journal of the ACM* 27(2):356–364, 1980.
11. Robert Nieuwenhuis and Albert Oliveras. Fast congruence closure and extensions. *Information and Computation* 205:557–580, 2007.
12. Robert E. Shostak. An algorithm for reasoning about equality. *Communications of the ACM* 21(7):583–585, 1978.
13. Aaron Stump, Clark W. Barrett, David L. Dill, and Jeremy Levitt. A decision procedure for an extensional theory of arrays. In Joseph Halpern (Ed.), *Proceedings of the 16th IEEE Symposium on Logic in Computer Science*, IEEE Computer Society Press, 2001.